


The problem



When reading from multiple sources, blocking on one of the sources could be bad.

An example of denial of service.

One solution: one process for every client. What are the pros and cons of this solution?

Another way to look at the problem

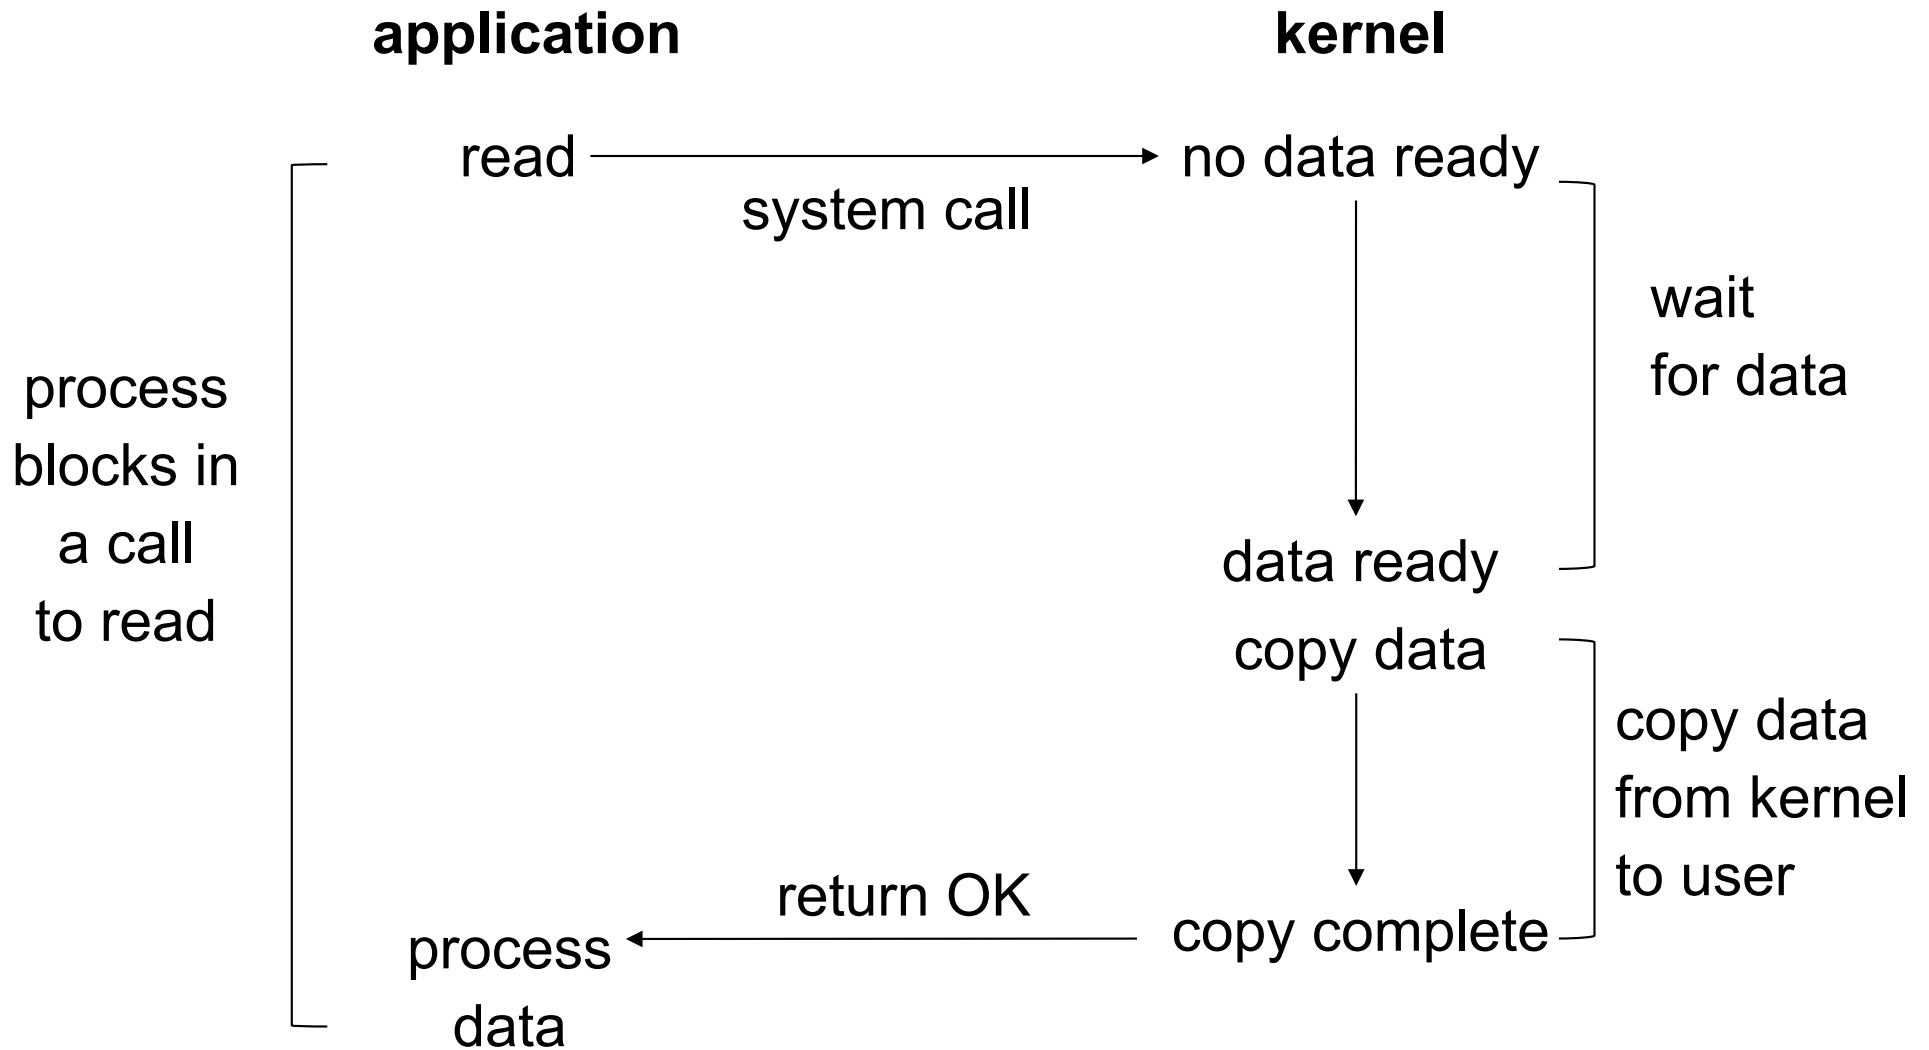
Server

```
while(1)
    accept a new connection
    for each existing connection
        read
        write
```

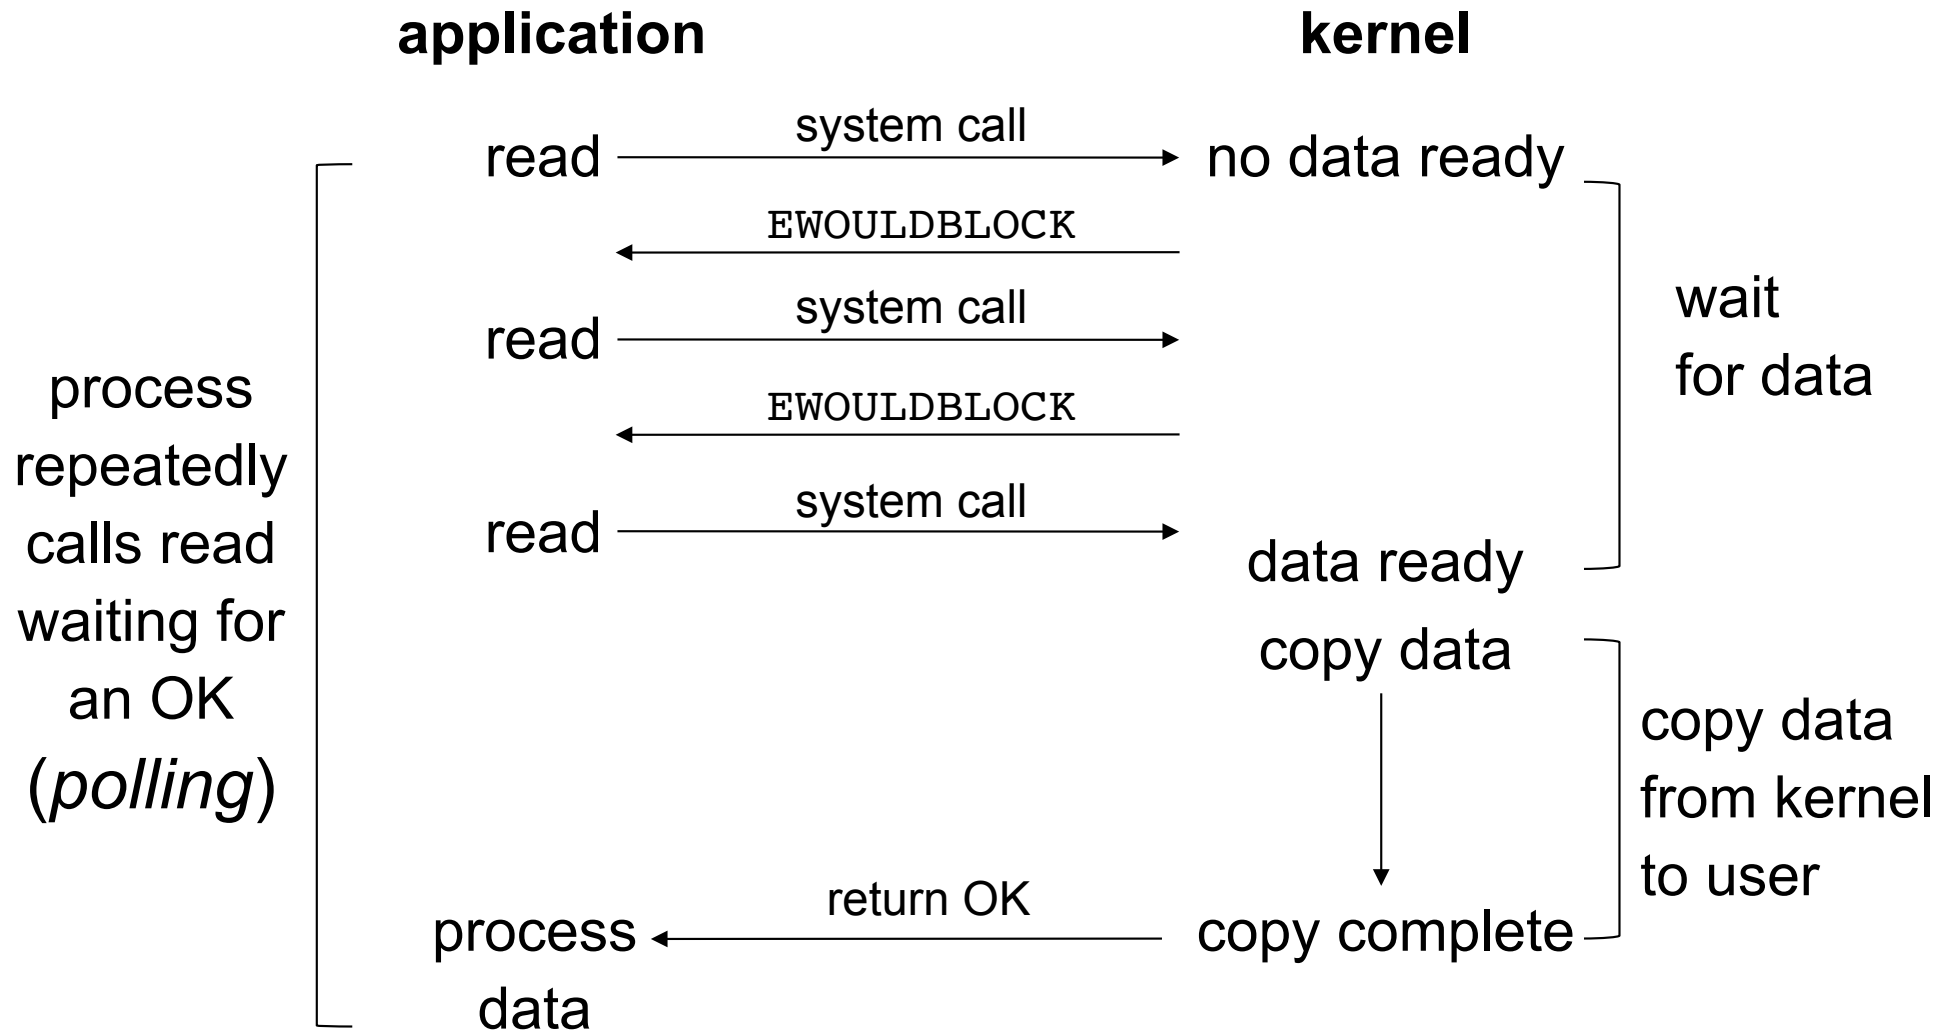
Which of the system calls might block indefinitely?
read and accept

So what happens if there is only one connection?

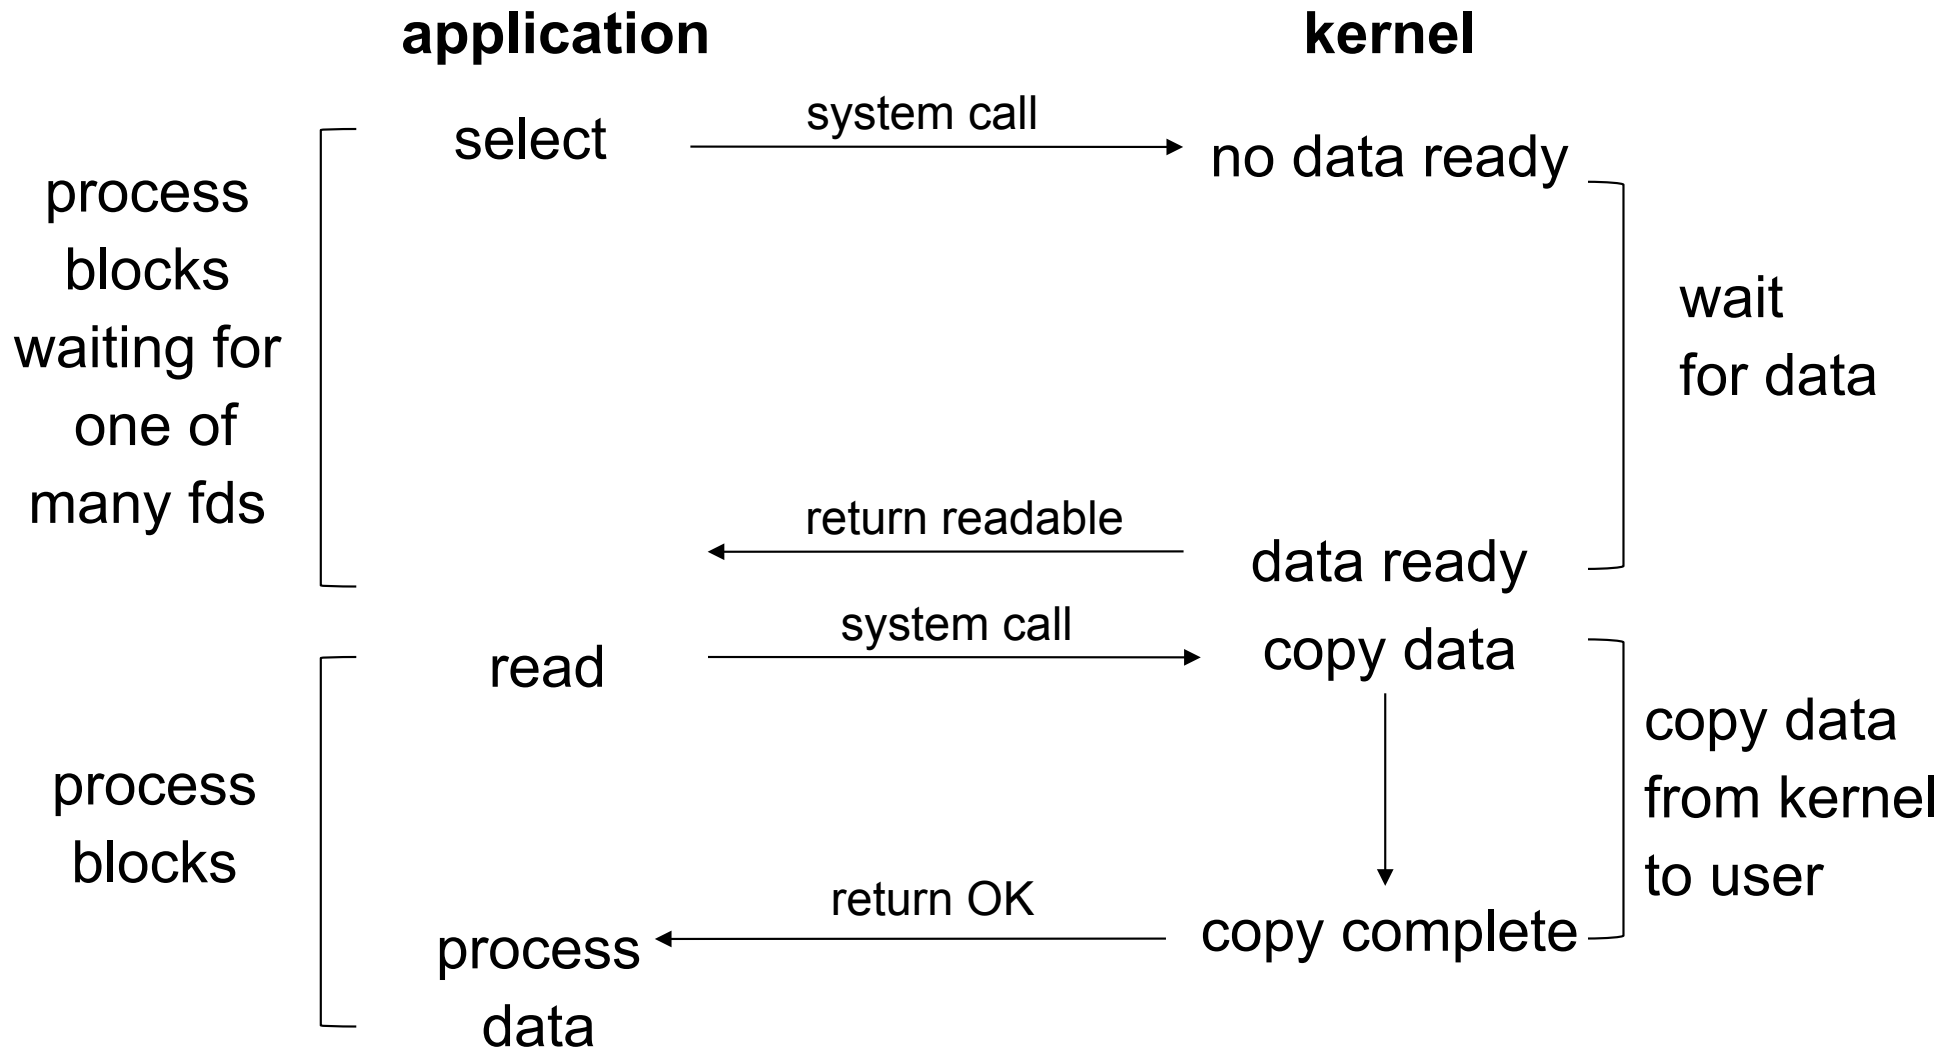
Blocking I/O Model



Nonblocking I/O Model



I/O Multiplexing Model



select()

```
int select(int maxfdp1,  
           fd_set *readset,  
           fd_set *writeset,  
           fd_set *exceptset,  
           const struct timeval *timeout);
```

A call to select returns when one of the file descriptors in one of the sets is ready for I/O.

If timeout is not NULL, then select returns when a descriptor is ready or timeout time has passed.

If timeout is 0, select returns immediately after checking descriptors.

Readiness

Ready to read when

- there is data in the receive buffer to be read

- end-of-file state on file descriptor

- the socket is a listening socket and there is a connection pending

- a socket error is pending

Ready to write when

- there is space available in the write buffer

- a socket error is pending

Exception condition pending when

- TCP out-of-band data

We are typically interested in when bytes are available to be read, but sometimes we use select on write or exception sets

select timeout

- The timeout specifies how long we're willing to wait for a fd to become ready

```
struct timeval {  
    long    tv_sec;        /* seconds */  
    long    tv_usec;       /* microseconds */  
};
```

- If timeout is NULL, wait forever (or until we catch a signal)
- If timeout is zero, test and return immediately
- Otherwise wait up to specified timeout
- `select` returns when a fd ready or we timeout

Descriptor sets

Typically implemented as an array of integers where each bit corresponds to a descriptor (except in Windows).


Implementation is hidden in the `fd_set` data type
`FD_SETSIZE` is the number of descriptors in the data type

`maxfdp1` specifies the number of descriptors to test

Macros:

```
void FD_ZERO(fd_set *fdset);
void FD_SET(int fd, fd_set *fdset);
void FD_CLR(int fd, fd_set *fdset);
int  FD_ISSET(int fd, fd_set *fdset);
```

Descriptor sets



	fd0	fd1	fd2	fd3	fd4	fd5	fd6
allset	0	0	0	1	1	0	1

$$\text{maxfd} + 1 = 7$$

After select:

	fd0	fd1	fd2	fd3	fd4	fd5	fd6
rset	0	0	0	1	0	0	0

select example

```
fd_set rfd;
struct timeval tv;
int retval;

FD_ZERO(&rfd); /* Watch stdin (fd 0) for input */
FD_SET(STDIN_FILENO, &rfd);
tv.tv_sec = 5; /* Wait up to five seconds. */
tv.tv_usec = 0;
retval = select(1, &rfd, NULL, NULL, &tv);
if (retval == -1)
    perror("select()");
else if (retval > 0)
    printf("Data is available now.\n");
    /* FD_ISSET(0, &rfd) will be true, can use read() */
else
    printf("No data within five seconds.\n");
```

```

for( ; ; ) {
    rset = allset;
    nready = Select(maxfd+1, &rset ,NULL,NULL,NULL);
    if(FD_ISSET(listenfd, &rset)) {
        connfd = Accept(listenfd, &caddr, &crlen);
        for(i = 0; i < FD_SETSIZE; i++)
            if(client[i] < 0) {
                client[i] = connfd; break;
            }
        FD_SET(connfd, &allset);
        if(connfd > maxfd) maxfd = connfd;
    }
    for(i = 0; i <= maxi; i++) {
        if(sockfd = client[i]) < 0) continue;
        if(FD_ISSET(sockfd, &rset))
            Read(sockfd, line, MAXLINE);
    }
}

```

```

for( ; ; ) {
    rset = allset;
    nready = Select(maxfd+1, &rset ,NULL,NULL,NULL);
    if(FD_ISSET(listenfd, &rset)) {
        connfd = Accept(listenfd, &caddr, &crlen);
        for(i = 0; i < FD_SETSIZE; i++)
            if(client[i] < 0) {
                client[i] = connfd; break;
            }
        FD_SET(connfd, &allset);
        if(connfd > maxfd) maxfd = connfd;
    }
    for(i = 0; i <= maxi; i++) {
        if(sockfd = client[i]) < 0) continue;
        if(FD_ISSET(sockfd, &rset))
            Read(sockfd, line, MAXLINE);
    }
}
}

```

```

for( ; ; ) {
    rset = allset;
    nready = Select(maxfd+1, &rset ,NULL,NULL,NULL);
    if(FD_ISSET(listenfd, &rset)) {
        connfd = Accept(listenfd, &caddr, &crlen);
        for(i = 0; i < FD_SETSIZE; i++)
            if(client[i] < 0) {
                client[i] = connfd; break;
            }
        FD_SET(connfd, &allset);
        if(connfd > maxfd) maxfd = connfd;
    }
    for(i = 0; i <= maxi; i++) {
        if(sockfd = client[i]) < 0) continue;
        if(FD_ISSET(sockfd, &rset))
            Read(sockfd, line, MAXLINE);
    }
}
}

```

End of Line


- There are two characters that determine end-of-line
 - Carriage return (CR, \r, ^M)
 - Line feed (LF, \n)
- Early operating systems defined their own conventions using one or both of CR and LF.
 - Unix: LF,
 - DOS/Windows: CR LF
 - Mac classic: CR

Network Line Ending


- Transferring data between machines with different operating systems, means deciding on a common line ending.
- CR LF is the standard
- (Of course it is possible with regular expression matching to mostly ignore this issue, but still better to conform.⁶)

Byte order

- **Big-endian**



- **Little-endian**



- Intel is little-endian, and Sparc is big-endian

Network byte order

- To communicate between machines with unknown or different “endian-ness” we convert numbers to network byte order (big-endian) before we send them.
- There are functions provided to do this:
 - `unsigned long htonl(unsigned long)`
 - `unsigned short htons(unsigned short)`
 - `unsigned long ntohl(unsigned long)`
 - `unsigned short ntohs(unsigned short)`

Arrays of bit strings

- FD_SETSIZE is bigger than 32.

```
struct bits {  
    unsigned int field[N];  
}  
typedef struct bits Bitstring;  
Bitstring a, b;  
setzero(&a);  
b = a;  
a.field[0] = ~0;
```

Setting and unsetting

```
int set(unsigned int bit, Bitstring *b) {  
    int index = bit / 32;  
    b->field[index] |= 1 << (bit % 32);  
    return 1;  
}
```

```
int unset(unsigned int bit, Bitstring *b) {  
    int index = bit / 32;  
    b->field[index] &= ~(1 << (bit % 32));  
}
```

Testing and emptying

```
int ifset(unsigned int bit, Bitstring *b) {  
    int index = bit / 32;  
    return ( (1 << (bit % 32))  
            & b->field[index]);  
}
```

```
int setzero(Bitstring *b){  
    if(memset(b,0, sizeof(Bitstring)) == NULL)  
        return 0;  
    else  
        return 1;  
}
```

Printing

```
char *intToBinary(unsigned int number) {  
    char *binaryString = malloc(32+1);  
    int i;  
    binaryString[32] = '\0';  
    for (i = 31; i >= 0; i--) {  
        binaryString[i] = ((number & 1) + '0');  
        number = number >> 1;  
    }  
    return binaryString;  
}
```